Name: Khushal Pareta

Student ID: 240840325031

#HIVE

Question 1:

1-> airports that are listed as both source and destination

```
SELECT ap.airport_id, ap.name
FROM airports ap
JOIN routes r1 ON ap.airport_id = r1.src_airport_id
JOIN routes r2 ON ap.airport_id = r2.dest_airport_id
LIMIT 10;
```

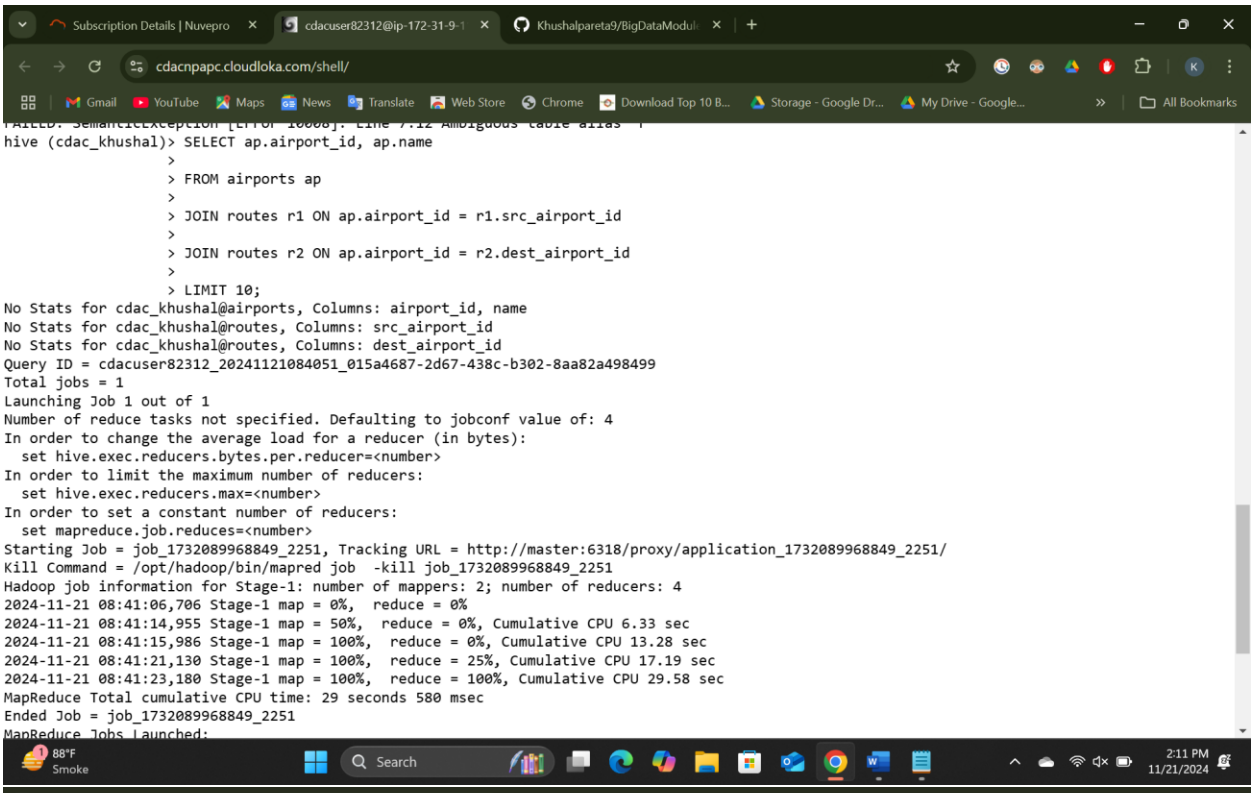

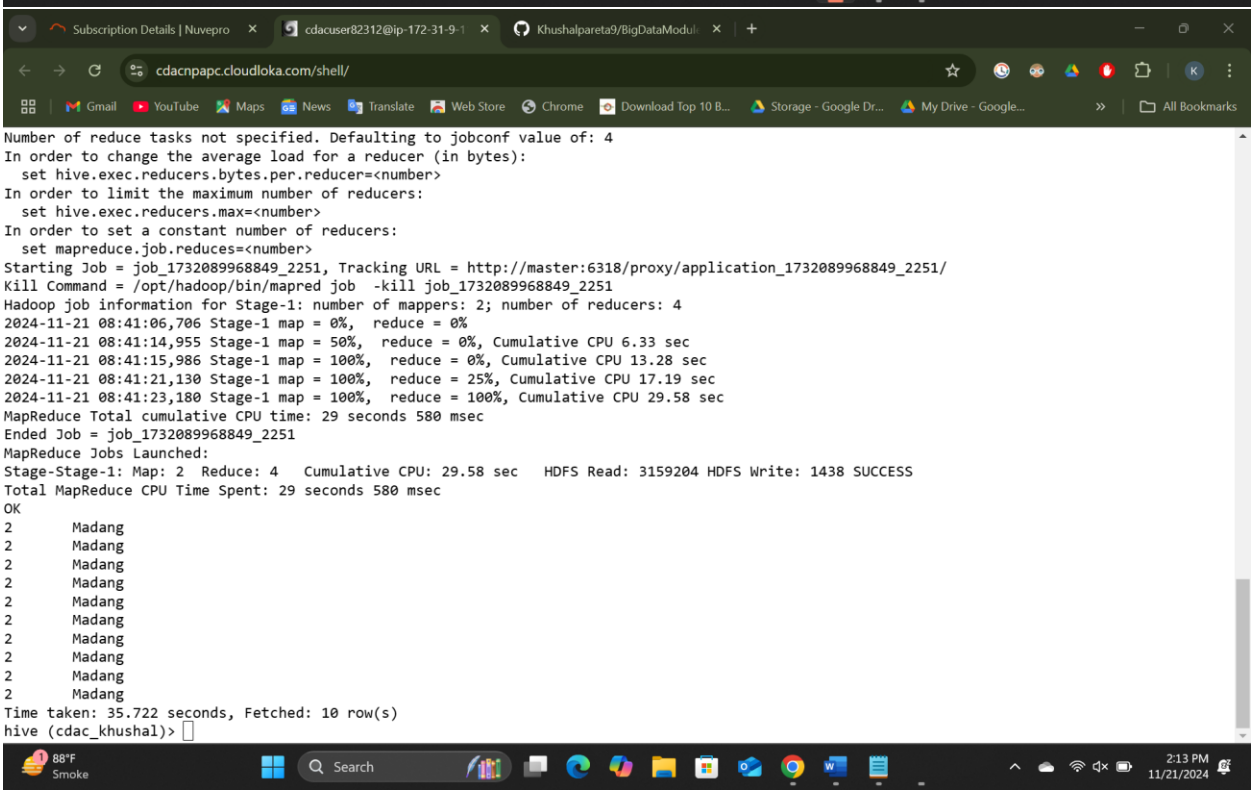

2 -> determine equipment that is used on highest number of routes

```
SELECT equipment, COUNT(equipment) AS HIGHEST_COUNT
FROM routes
GROUP BY equipment, src_airport_id, dest_airport_id
ORDER BY HIGHEST_COUNT DESC
LIMIT 1;
```

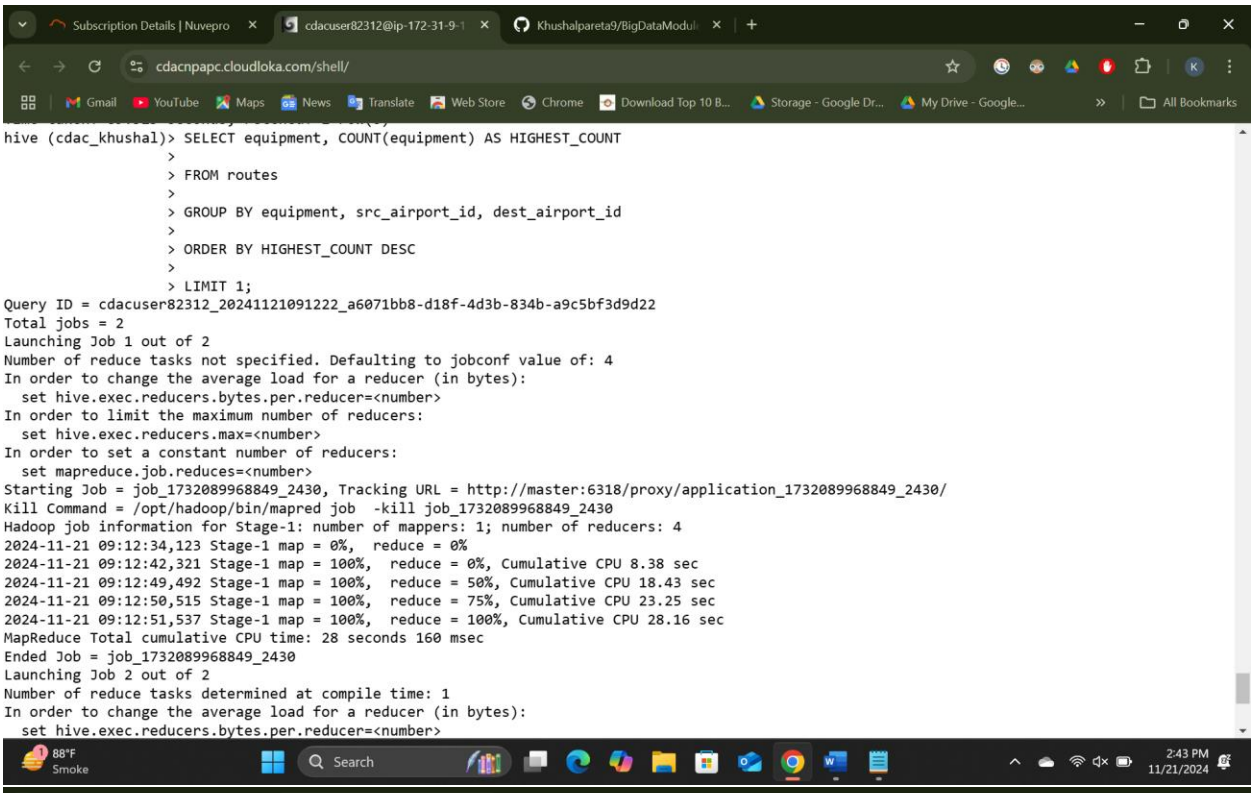

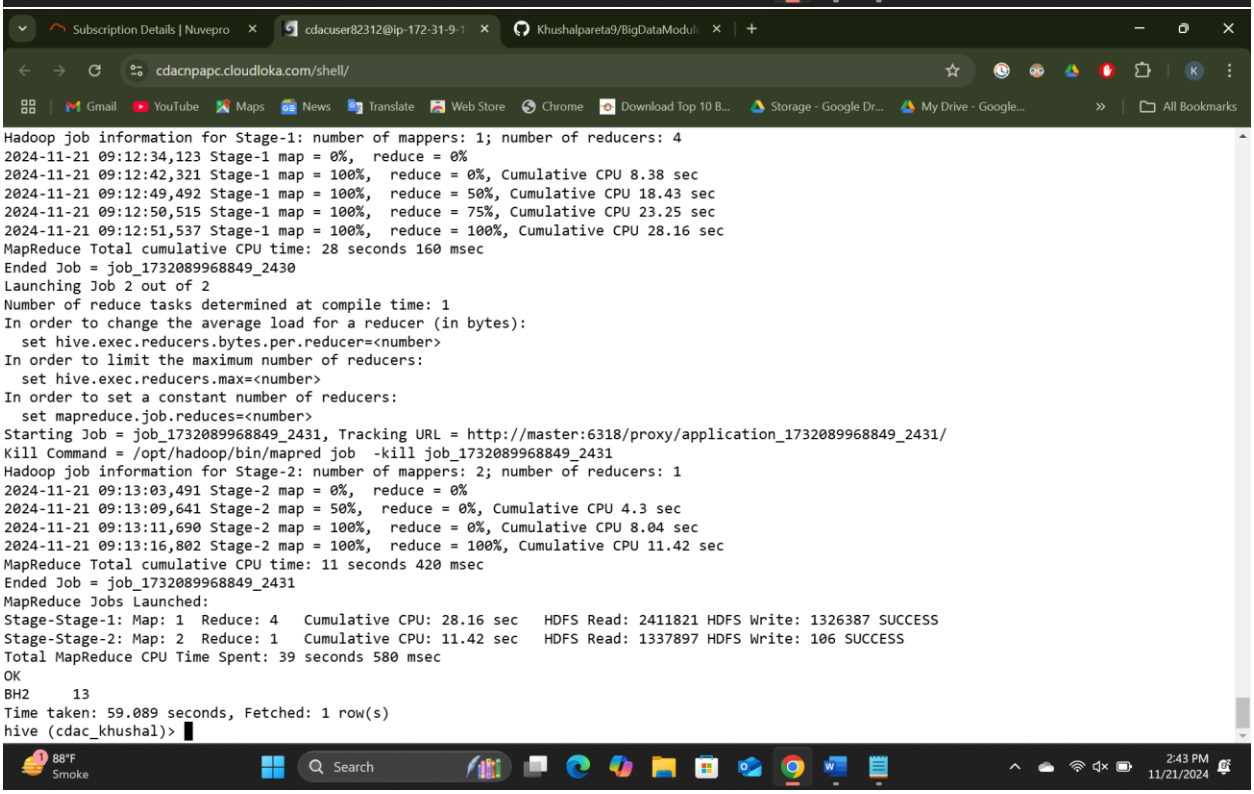

3 - > Airline which operates the highest number of routes and count of those routes

```
SELECT a.name, COUNT(a.airline_id) AS ROUTE_COUNT
FROM airlines a
JOIN routes r ON a.airline_id = r.airline_id
GROUP BY a.name, r.src_airport_id, r.dest_airport_id
ORDER BY ROUTE_COUNT DESC
LIMIT 1;
```

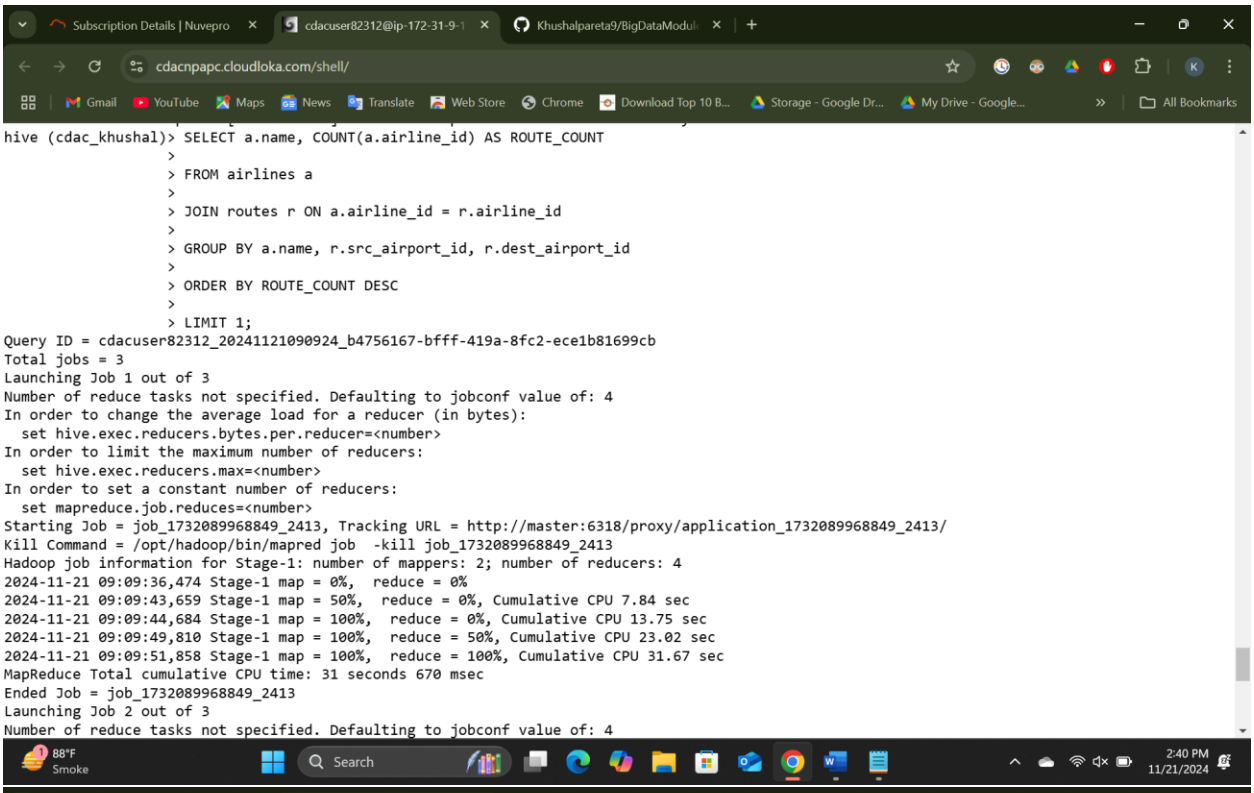
Subscription Details | Nuvepro
cdacuser82312@ip-172-31-9-1
Khushalpareta9/BigDataModul
cdacnpapc.cloudloka.com/shell/
Gmail
YouTube
Maps
News
Translate
Web Store
Chrome
Download Top 10 B...
Storage - Google Dr...
My Drive - Google...
All Bookmarks
hive (cdac_khushal)> SELECT a.name, COUNT(a.airline_id) AS ROUTE_COUNT
> FROM airlines a
> JOIN routes r ON a.airline_id = r.airline_id
> GROUP BY a.name, r.src_airport_id, r.dest_airport_id
> ORDER BY ROUTE_COUNT DESC
> LIMIT 1;
Query ID = cdacuser82312_20241121090924_b4756167-bfff-419a-8fc2-ece1b81699cb
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks not specified. Defaulting to jobconf value of: 4
In order to change the average load for a reducer (in bytes):
set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
set mapreduce.job.reduces=<number>
Starting Job = job_1732089968849_2413, Tracking URL = http://master:6318/proxy/application_1732089968849_2413/
Kill Command = /opt/hadoop/bin/mapred job -kill job_1732089968849_2413
Hadoop job information for Stage-1: number of mappers: 2; number of reducers: 4
2024-11-21 09:09:36,474 Stage-1 map = 0%, reduce = 0%
2024-11-21 09:09:43,659 Stage-1 map = 50%, reduce = 0%, Cumulative CPU 7.84 sec
2024-11-21 09:09:44,684 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 13.75 sec
2024-11-21 09:09:49,810 Stage-1 map = 100%, reduce = 50%, Cumulative CPU 23.02 sec
2024-11-21 09:09:51,858 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 31.67 sec
MapReduce Total cumulative CPU time: 31 seconds 670 msec
Ended Job = job_1732089968849_2413
Launching Job 2 out of 3
Number of reduce tasks not specified. Defaulting to jobconf value of: 4
88°F
Smoke
Search
2:40 PM
11/21/2024

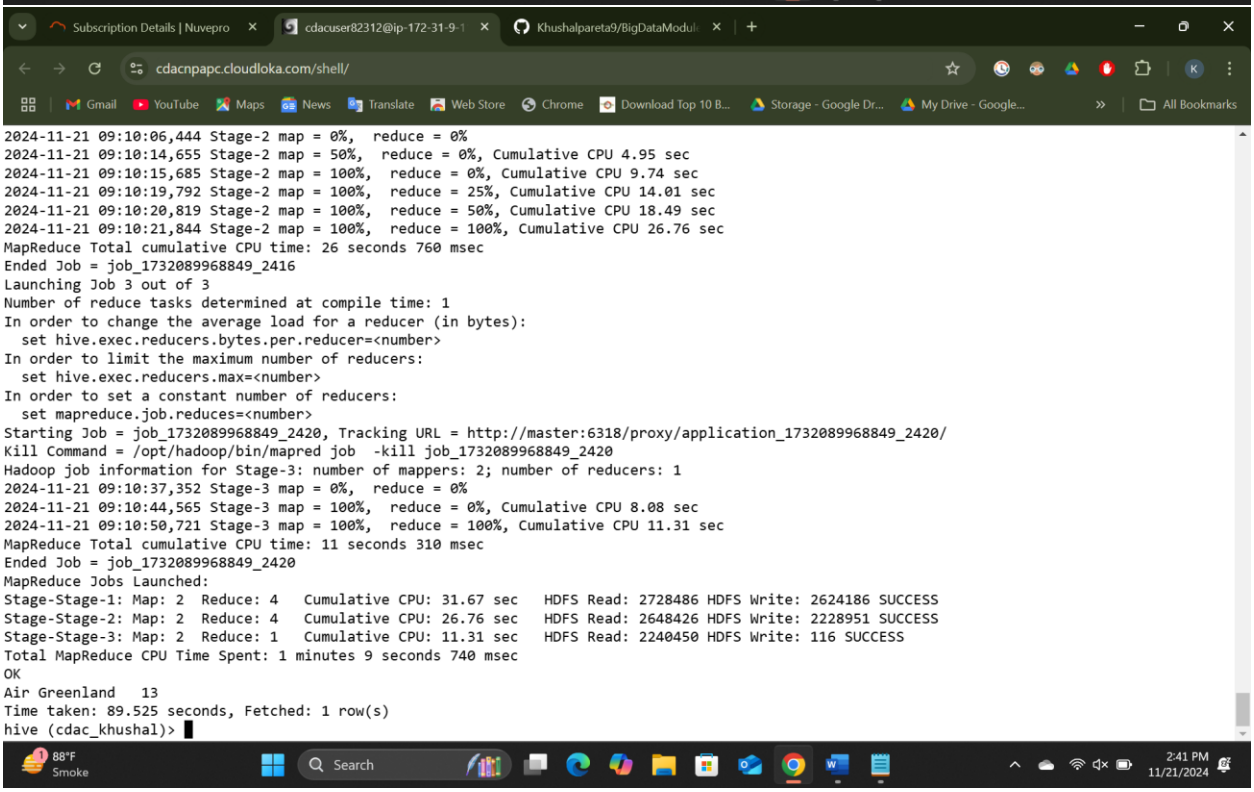
Subscription Details | Nuvepro
cdacuser82312@ip-172-31-9-1
Khushalpareta9/BigDataModul
cdacnpapc.cloudloka.com/shell/
Gmail
YouTube
Maps
News
Translate
Web Store
Chrome
Download Top 10 B...
Storage - Google Dr...
My Drive - Google...
All Bookmarks
2024-11-21 09:10:06,444 Stage-2 map = 0%, reduce = 0%
2024-11-21 09:10:14,655 Stage-2 map = 50%, reduce = 0%, Cumulative CPU 4.95 sec
2024-11-21 09:10:15,685 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 9.74 sec
2024-11-21 09:10:19,792 Stage-2 map = 100%, reduce = 25%, Cumulative CPU 14.01 sec
2024-11-21 09:10:20,819 Stage-2 map = 100%, reduce = 50%, Cumulative CPU 18.49 sec
2024-11-21 09:10:21,844 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 26.76 sec
MapReduce Total cumulative CPU time: 26 seconds 760 msec
Ended Job = job_1732089968849_2416
Launching Job 3 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
set mapreduce.job.reduces=<number>
Starting Job = job_1732089968849_2420, Tracking URL = http://master:6318/proxy/application_1732089968849_2420/
Kill Command = /opt/hadoop/bin/mapred job -kill job_1732089968849_2420
Hadoop job information for Stage-3: number of mappers: 2; number of reducers: 1
2024-11-21 09:10:37,352 Stage-3 map = 0%, reduce = 0%
2024-11-21 09:10:44,565 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 8.08 sec
2024-11-21 09:10:50,721 Stage-3 map = 100%, reduce = 100%, Cumulative CPU 11.31 sec
MapReduce Total cumulative CPU time: 11 seconds 310 msec
Ended Job = job_1732089968849_2420
MapReduce Jobs Launched:
Stage-Stage-1: Map: 2 Reduce: 4 Cumulative CPU: 31.67 sec HDFS Read: 2728486 HDFS Write: 2624186 SUCCESS
Stage-Stage-2: Map: 2 Reduce: 4 Cumulative CPU: 26.76 sec HDFS Read: 2648426 HDFS Write: 2228951 SUCCESS
Stage-Stage-3: Map: 2 Reduce: 1 Cumulative CPU: 11.31 sec HDFS Read: 2240450 HDFS Write: 116 SUCCESS
Total MapReduce CPU Time Spent: 1 minutes 9 seconds 740 msec
OK
Air Greenland 13
Time taken: 89.525 seconds, Fetched: 1 row(s)
hive (cdac_khushal)>
88°F
Smoke
Search
2:41 PM
11/21/2024

Question 2:

1 -> Create a partition table for the source_airport, write a sql query to create this table and insert data into it.

SET hive.exec.dynamic.partition=true;

SET hive.exec.dynamic.partiton.mode=nonstrict;

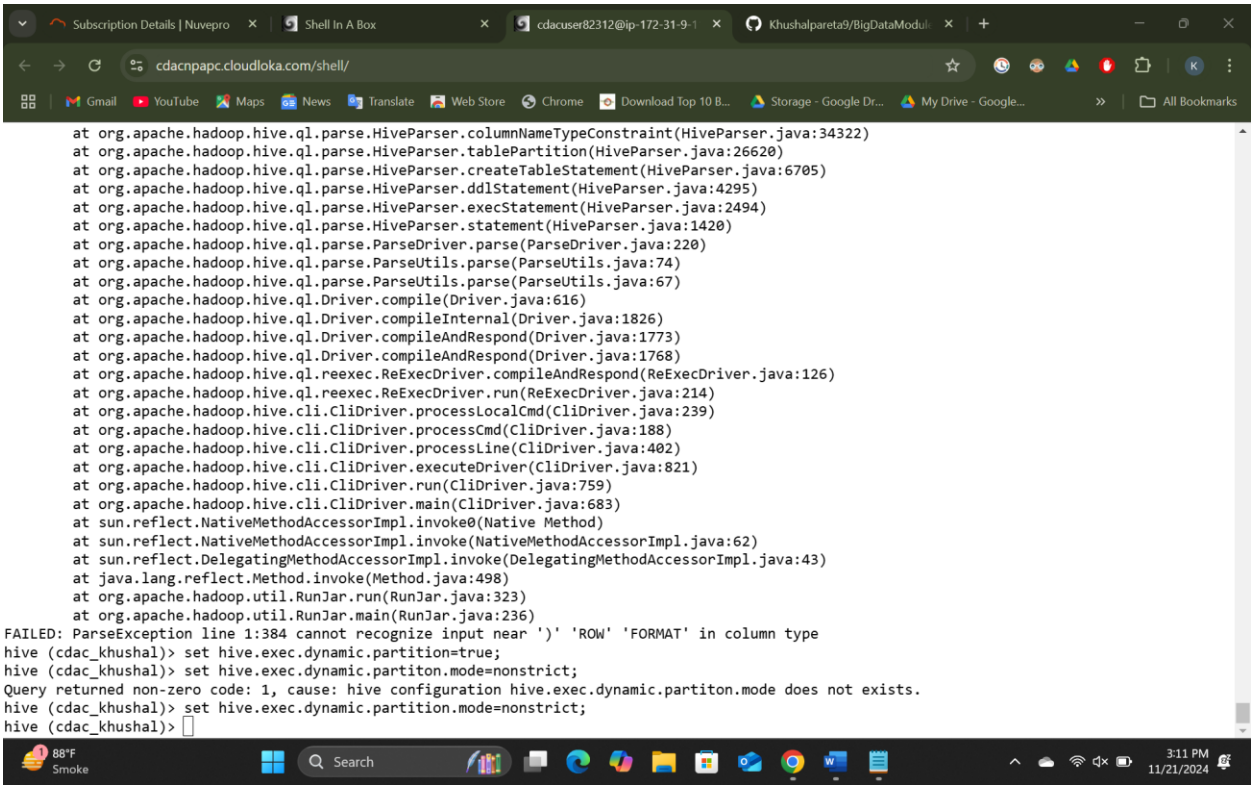

CREATE TABLE partitioned_table (

airline_iata string, airline_id int, src_airport_id int, dest_airport_iata string, dest_airport_id int, codeshare string, stops int, equipment string

)

PARTITIONED BY (source_airport_iata='JFK')

ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE;

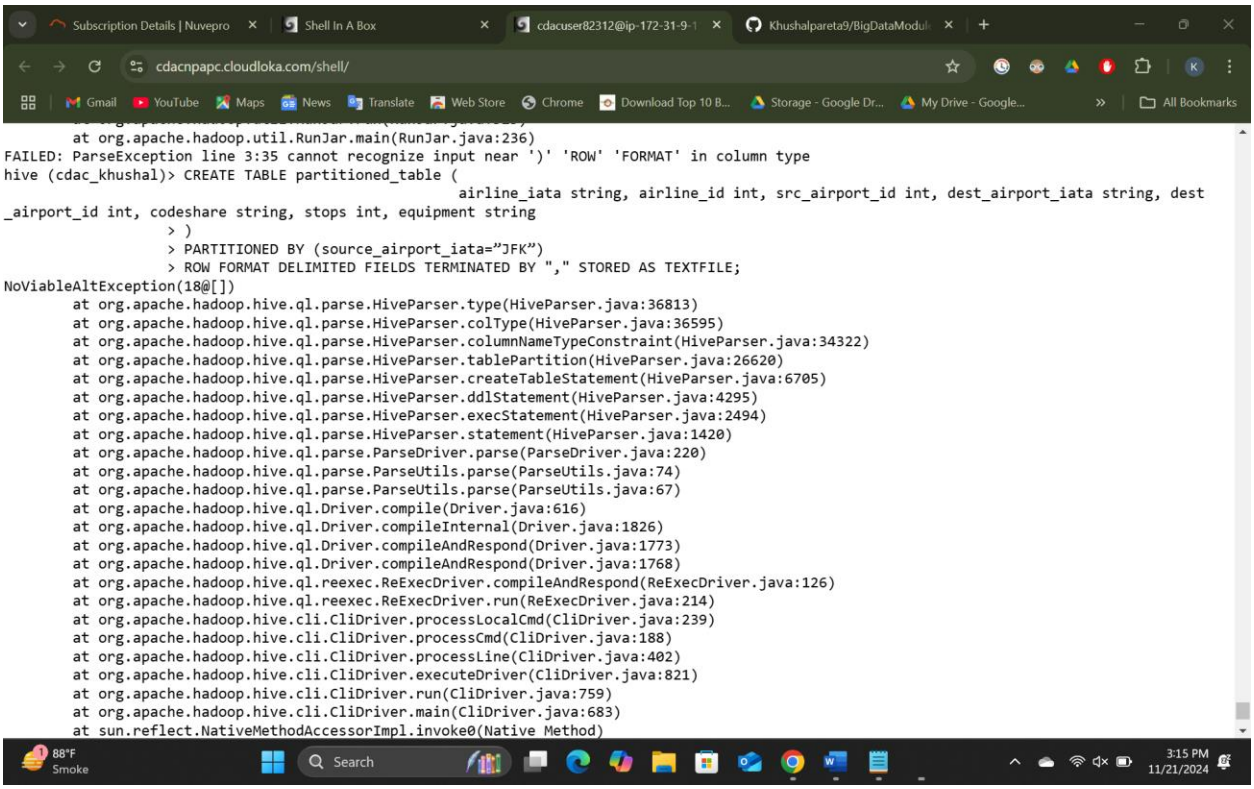

2 ->

INSERT INTO TABLE partitioned_table SELECT * FROM airports WHERE src_airport_iata='JFK'

3 ->

SELECT * FROM partitioned_table WHERE src_airport_iata='LAX';

4 ->

SHOW paritions partitioned_table

#SPARK

Question 1:

df = spark.read.csv('/user/cdacuser82312/sparkAirlines/spark_airlines.csv', header=True, inferSchema=True)

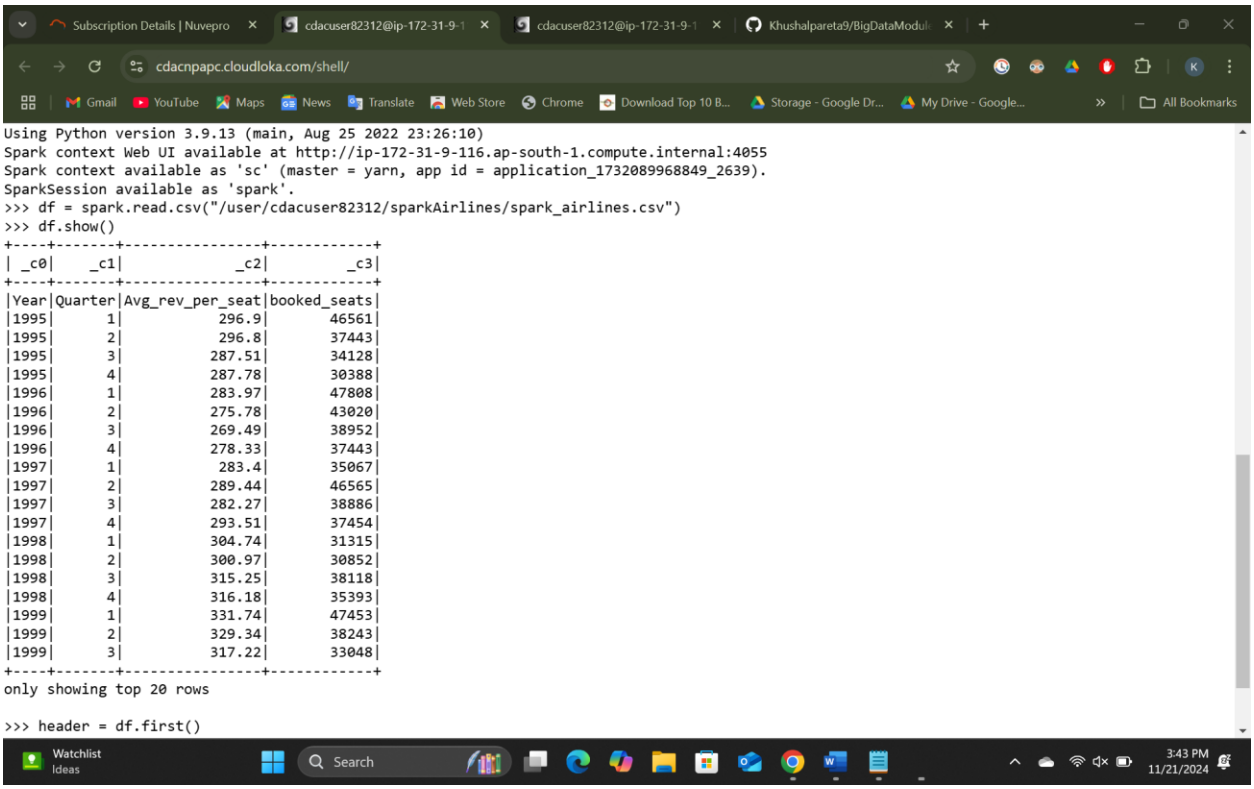

1 ->

From pyspark.sql.functions import sum

df.groupBy('Year', 'Quarter').agg(count(sum('booked_seats')) > 40000).show()

2 ->

df.groupBy('Year').show()

QUESTION 2:

1->

From pyspark.sql.functions import sum, avg, min

df.groupBy('Year', 'Quarter').agg(sum('Avg_rev_per_seat').alias('TotalRevenuePerSeat'),
avg('Avg_rev_per_seat').alias('AverageRevenuePerSeat'),
min('Avg_rev_per_seat').alias('MinimumRevenuePerSeat')).show()

2 ->

df.groupBy('Year', 'Quarter').agg(count(avg_rev_per_seat).alias('Total_Count') > 290)).show()

3 ->

From pyspark.sql.functions import sum

df.groupBy('Year').agg(sum('booked_seats').alias('Total_Booked_Seats)).show()

4 ->

df.groupBy('year').show()

5 ->

From pyspark.sql.functions import sum

df.groupBy('Year').agg(sum('avg_rev_per_seat').alias('Total_Avg_Revenue')).show()